Exact and Error-Bounded VUS Evaluation through Sparse Events and Width Geometry

Working manuscript; unsubmitted

Author metadata pending
Affiliation pending
Country pending

ABSTRACT

Volume Under the Surface (VUS) averages time-series anomaly-detection areas over score thresholds and tolerance widths. We reduce computation along both axes while preserving a pinned metric definition. First, thresholds outside the maximum anomaly support contribute analytically integrable ROC rectangles and zero PR area. Sorting support scores and counting remaining scores yields exact scalar VUS-ROC and VUS-PR in $O(n \log(Q+1) + (L+1)Q)$ time and $O(n+Q)$ memory, where Q is the expanded support size and L the largest width. Support restriction and cumulative sums are prior techniques; our reduction eliminates inactive threshold events. Second, expanded-interval merge events partition the width range into regions with ordered ROC coordinates. A probability-box expectation bound encloses PR without assuming monotone PR area. Adaptive refinement bounds the discrete width average to a requested half-width in real arithmetic; the float64 implementation is numerically audited, not a directed-rounding certificate. Experiments use 350 public series and 1050 released score vectors. Against released fast-VUS computing PR only, the paired ROC/PR evaluator has median speedups of 8.38× and 4.55× on two CPU systems. Against the fastest tested comparator, including dense methods, the medians are 2.04× and 1.35×; Sparse VUS is slower or equal in 209 and 574 of 2,100 cases, respectively. Width refinement evaluates fewer curves but is often slower at tight tolerances. Incremental PR racing gives median speedups at width 300, yet increases the sum of per-task median latencies on both machines; both extensions remain optional. The reference's 250-threshold approximation changes 84 of 2,100 PR pair orders; 75 and 74 persist after two fixed padding interventions. The results distinguish exact-integration gains from the cost of error-bounded width reduction.

PVLDB Reference Format:
Author metadata pending. Exact and Error-Bounded VUS Evaluation through Sparse Events and Width Geometry. PVLDB, 20(1): XXX-XXX, 2027.
doi:XX.XX/XXX.XX

PVLDB Artifact Availability:
The source code, data, and/or other artifacts have been made available at https://github.com/mandu5/sparse-vus/releases/tag/v0.2.0.

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit https://creativecommons.org/licenses/by-nc-nd/4.0/ to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.
Proceedings of the VLDB Endowment, Vol. 20, No. 1 ISSN 2150-8097.
doi:XX.XX/XXX.XX

1 INTRODUCTION

VUS averages range-aware ROC or precision–recall integrals across tolerance widths, accounting for both overlap and whether an anomaly interval is hit [1, 6]. Efficient evaluation matters: ESA-ADB explicitly excluded VUS because of computational cost [3]. That observation motivates speed improvements; it does not establish suitability for ESA-ADB's other evaluation requirements.

The official optimized evaluator partitions static and dynamic labels and restricts work to maximum-width anomaly support [1]. Klein's released fast-VUS further sorts support scores and uses cumulative weighted sums [2]. These optimizations already avoid much repeated work. We instead ask which threshold events can be eliminated altogether, and which widths can be omitted with a quantified error. Only positions near labeled anomalies change the weighted positive state. Other thresholds have an analytically integrable effect, but removing them requires preserving a ROC rectangle before each support transition.

Width reduction requires a different argument. Expanded anomaly intervals can merge, changing existence reward discontinuously. Between merge events, we prove coordinate order and monotone ROC area. PR area need not be monotone even there. We express it as an expectation under recall increments and apply established p-box bounds [9, 10]. This yields an adaptive discrete-width average with separate lower and upper limits, rather than a fixed width subsample with unknown error.

Exact evaluation also avoids an accuracy limitation of a fixed uniform-rank threshold grid. At width zero, place the only anomaly at score rank two and append increasingly many lower-scored negatives. Exact PR remains $1/2$, whereas a fixed T-threshold grid first detects it at rank $1 + \lfloor (n-1)/(T-1) \rfloor$, and its PR tends to zero. Appendix A gives a version whose absolute error approaches one for any fixed width range. This elementary counterexample motivates retaining all thresholds; it is not specific to the sparse evaluator.

Our target is the pair of scalar outputs of `RangeAUC_volume_opt`, with `thre=n`, in TheDatumOrg/VUS commit `467bace0` (the full revision and file hashes are in the artifact). Tied repeated thresholds add zero area. The target includes dynamic label mass, sum-and-clip overlapping ramps, expanded-interval existence reward, ROC trapezoids, the terminal $(1, 1)$ ROC point, and the PR right-endpoint step integral. Other historical VUS routines can use different semantics. We return scalar integrals, not the full surface coordinates.

2 EXACT TARGET AND SPARSE EVENTS

Let $y_i \in \{0,1\}$ contain both classes and let s_i be a finite score, for $i = 1, \dots, n$. Larger scores indicate greater anomaly evidence. Write $P_0 = \sum_i y_i$ and let $w \in \{0, \dots, L\}$ be a tolerance width. The official continuous label $a_i^{(w)}$ equals one on original anomaly intervals. Each interval contributes $\sqrt{1 - d/w}$ outside its endpoints for integer distance $1 \le d \le \lfloor w/2 \rfloor$; contributions are summed and capped at one. For $w = 0$ no ramp is added.

At threshold t, define $p_i(t) = \mathbf{1}\{s_i \ge t\}$ and

$$C(t) = \sum_i p_i(t), \qquad \mathrm{TP}(t) = \sum_i a_i^{(w)} p_i(t), \tag{1}$$

$$B(t) = \sum_i (1 - y_i) a_i^{(w)} p_i(t), \tag{2}$$

$$P(t) = P_0 + B(t)/2, \qquad N(t) = n - P(t). \tag{3}$$

The identity for $P(t)$ follows by resetting original positive labels to one after thresholding the buffer labels, then averaging their mass with P_0. Let K_w be the number of expanded intervals and $H(t)$ the number containing a predicted positive. Overlapping expansions merge; adjacent nonoverlapping expansions remain separate, matching the reference. The curve coordinates are

$$\mathrm{TPR}(t) = \min\{\mathrm{TP}(t)/P(t), 1\} \frac{H(t)}{K_w}, \tag{4}$$

$$\mathrm{FPR}(t) = \frac{C(t) - \mathrm{TP}(t)}{N(t)}, \qquad \mathrm{Prec}(t) = \frac{\mathrm{TP}(t)}{C(t)}. \tag{5}$$

Here $N(t) > 0$ because $P(t) \le (n + P_0)/2 < n$. ROC uses trapezoids from the origin through all distinct descending thresholds to the explicit terminal $(1, 1)$. PR is $\sum_j (\mathrm{TPR}_j - \mathrm{TPR}_{j-1}) \mathrm{Prec}_j$, with initial recall zero. VUS is the arithmetic mean of these areas over $w = 0, \dots, L$.

Let S_L be the union of all positions in maximum-width expanded intervals, $Q = |S_L|$, and let $G \le Q$ be the number of distinct scores in S_L. The distinction matters: Q counts positions, not distinct thresholds. All nonzero weights and all positions capable of changing existence reward lie in S_L, for every evaluated width.

Lemma 1 (Negative-only runs). *Between consecutive distinct support scores, the states* TP*, B, P, N, H, and* TPR *are constant. Intermediate thresholds contribute only a horizontal ROC rectangle and zero PR area.*

Proof. Every newly admitted point outside S_L has zero weight and belongs to no expanded interval at width $w \le L$. Thus only C changes. Since $\mathrm{FPR} = (C - \mathrm{TP})/N$ with fixed TP, N, all intermediate ROC points have the same height. Trapezoid widths telescope to the endpoint difference. Every PR recall increment is zero. This includes arbitrary ties among off-support scores. □

For a descending support event j at score v_j, compute $c_j^{>} = |\{i : s_i > v_j\}|$ and $c_j^{\ge} = |\{i : s_i \ge v_j\}|$. Write $T = \mathrm{TP}$, $R = \mathrm{TPR}$, and $F = \mathrm{FPR}$ for compact event notation. Suppose the preceding event leaves state (T_o, N_o, R_o, c_o) and the new event has coordinates $(F_j, R_j, \mathrm{Prec}_j)$. Its exact contribution is

$$f_j^{-} = (c_j^{>} - T_o)/N_o, \tag{6}$$

$$A_j^{\mathrm{ROC}} = (c_j^{>} - c_o) R_o / N_o + (F_j - f_j^{-})(R_j + R_o)/2, \tag{7}$$

$$A_j^{\mathrm{PR}} = (R_j - R_o) \mathrm{Prec}_j. \tag{8}$$

The initial state is $(0, n - P_0, 0, 0)$; thereafter $c_o = c_{j-1}^{\ge}$. The first term is the negative rectangle; the second admits the whole tied support event, including off-support points tied with it. Omitting the rectangle and connecting retained event endpoints directly is generally wrong. After the last event, the remaining negative tail and terminal ROC transition contribute

$$\frac{(n - c_G^{\ge}) R_G}{N_G} + \left(1 - \frac{n - T_G}{N_G}\right) \frac{1 + R_G}{2}. \tag{9}$$

For example, $y = (1, 0, 1, 0)$, $s = (3, 2, 1, 0)$, and $L = 0$ give exact ROC area 0.625. Connecting only support endpoints incorrectly gives 0.8125. The exact step-PR area is 0.75.

Lemma 2 (Two interval distances). *For a non-original position, the official buffer weight is determined by its two nearest distinct anomaly intervals.*

Proof. Each active ramp contributes at least $\sqrt{1/2}$. If at least two intervals contribute, their sum exceeds one and clipping returns one. Otherwise the unique contribution is determined by the nearest interval. Hence the weight is zero when $d_1 > \lfloor w/2 \rfloor$, one when $d_2 \le \lfloor w/2 \rfloor$, and $\sqrt{1 - d_1/w}$ otherwise. Original positive positions remain one. The nearest two distances are computed from the two neighboring intervals on either side in sorted interval order. □

Theorem 1 (Support-sensitive exact evaluation). *The two pinned scalar VUS integrals can be computed in $O(n \log(Q + 1) + (L + 1)Q)$ time and $O(n + Q)$ memory in a unit-cost real-arithmetic model with score comparisons and square roots. The returned values equal the reference integrals in exact arithmetic.*

Proof. Extract intervals and S_L in $O(n + Q)$ time. Sort the Q support scores and group their G distinct values. Binary-search each of the n scores among these values and accumulate counts to obtain every $c_j^{>}, c_j^{\ge}$, without sorting off-support scores. This costs $O(Q \log Q + n \log(G + 1))$, bounded by $O(n \log(Q + 1))$ because $1 \le G \le Q \le n$. Precomputing the two nearest interval distances costs $O(Q \log(K + 1))$, absorbed by the same bound since the number of original intervals $K \le Q$.

For each width, construct weights in $O(Q)$ using Lemma 2. Assign disjoint expanded interval memberships in $O(Q + K)$ and find each interval's earliest support-score group. Grouped weighted sums and cumulative sums yield TP, B, and H for all G events in $O(Q + G + K) = O(Q)$ time. Lemma 1 and the event formulas reproduce every omitted threshold contribution, including the terminal tail. Average over $L + 1$ widths. Arrays for the input, a position map, support positions, grouped states, and interval metadata require $O(n + Q)$ memory. Floating-point implementations approximate this exact-arithmetic identity; their numerical error is measured separately. □

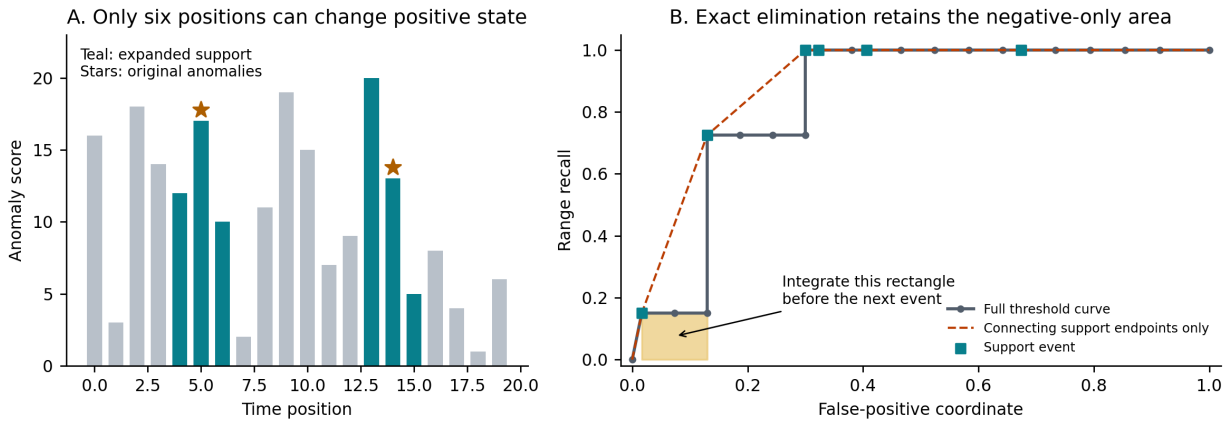


Figure 1: A width-two example. Negative-only thresholds move horizontally along the ROC curve. Integrating these rectangles and the subsequent tied support transitions preserves the full-threshold area (0.841642); joining support endpoints alone gives 0.897580.

The compiled scalar kernel streams the width average. The NumPy diagnostic implementation retains per-width areas and uses an additional $O(L)$ space. For dense support, $Q = \Theta(n)$, no asymptotic improvement over a dense sweep is promised. The analysis establishes an upper bound.

3 ERROR-BOUNDED REFINEMENT ALONG THE WIDTH AXIS

Exact threshold integration leaves a second cost: evaluating every integer width. We next bound the scalar average while omitting selected widths. This is a different reduction from restricting thresholds. It uses the geometry of expanded anomaly intervals and a probability-box (p-box) expectation bound; the latter is established prior work [9, 10].

3.1 Where width monotonicity holds

For consecutive original intervals ending at e_i and starting at s_{i+1}, their expansions first overlap at

$$c_i = 2\left\lceil \frac{s_{i+1} - e_i}{2} \right\rceil.$$

The strict overlap convention matters: merely adjacent expansions do not merge. Partition $0, \dots, L$ immediately before every distinct $c_i \le L$. Within each resulting integer interval $[a, b]$, the number K of expanded segments is constant, each segment expands without merging, and each threshold's hit count is nondecreasing.

Lemma 3 (Coordinate order within a merge-free interval). *For a fixed threshold and $a \le w \le b$ in one such interval,*

$$R_a \le R_w \le R_b, \quad \mathrm{Prec}_a \le \mathrm{Prec}_w \le \mathrm{Prec}_b, \quad F_a \ge F_w \ge F_b.$$

At any fixed width, recall and FPR are nondecreasing as the threshold falls.

Proof. Let O be the number of admitted original positives and B_w the admitted buffer mass. Both O and the prediction count C are fixed across widths, while B_w is nondecreasing. The unclipped overlap recall is $(O + B_w)/(P_0 + B_w/2)$, whose derivative in B_w is $(P_0 - O/2)/(P_0 + B_w/2)^2 > 0$. Multiplication by the nondecreasing existence reward H_w/K and clipping preserve order. Precision is $(O + B_w)/C$. For FPR,

$$\frac{\partial F}{\partial B} = \frac{-(n - P_0) + (C - O)/2}{(n - P_0 - B/2)^2} < 0,$$

since $C - O \le n - P_0$. Across descending thresholds, original-positive admissions increase O; other admissions increase B and increase the false-positive numerator by $1 - a_i^{(w)} \ge 0$ while decreasing its denominator. Thus both curve coordinates are nondecreasing in threshold order. Ties enter together. □

Theorem 2 (ROC width enclosure). *Within a merge-free interval, the exact ROC area is nondecreasing in width. Consequently every omitted width has area in* $[\mathrm{ROC}_a, \mathrm{ROC}_b]$.

Proof. Use the common full threshold sequence, with the prescribed origin and terminal point. Its trapezoidal area is coordinate-wise nondecreasing in each recall and nonincreasing in each interior FPR:

$$\frac{\partial A}{\partial R_j} = \frac{F_{j+1} - F_{j-1}}{2} \ge 0, \qquad \frac{\partial A}{\partial F_j} = \frac{R_{j-1} - R_{j+1}}{2} \le 0.$$

Linear interpolation between the endpoint coordinate sequences preserves their threshold monotonicity. Apply the coordinate-order lemma along this interpolation. Sparse event integration gives the same scalar area. □

Neither global ROC monotonicity nor PR-area monotonicity is assumed. For example, with $n = 20$, positive positions $\{0, 2, 19\}$ and scores $(20, 4, 19, 3, 18, 17, 16, 15, 14, 13, 12, 11, 10, 9, 8, 7, 6, 5, 2, 1)$, ROC drops from 0.444444 at width one to 0.405557 at width two when three expanded segments become two. Even without a merge, scalar endpoint PR values need not enclose the interior. Take $n = 700$, positive positions $100, \dots, 199$ and 600, and initially zero scores. Assign score three to positions $100, \dots, 149$, score one

to 150, … , 199 and 600, and score two to 0, … , 99, 200, … , 549, and 599. The first merge occurs at width 402. Direct evaluation gives PR 0.385206 at width one, 0.366313 at width two, and 0.909134 at width 400. Thus the width-two area lies below both endpoint values on $[1, 400]$. The artifact verifies this witness independently.

3.2 PR bounds without monotone precision

Index the common descending support events by $j = 1, \dots, G$. At every width, $q_j = R_j - R_{j-1} \geq 0$ and $\sum_j q_j = 1$, since the last support event has recall one. Hence PR is an expectation $\sum_j q_j \mathrm{Prec}_j$, with cumulative distribution R_j. For $z \in (0, 1)$, define the feasible quantile window

$$\ell(z) = \min\{j : R_{b,j} \geq z\}, \qquad u(z) = \min\{j : R_{a,j} \geq z\}.$$

The coordinate-order lemma implies that the quantile index for any intermediate width belongs to $[\ell(z), u(z)]$.

Theorem 3 (PR width enclosure). *Every width in $[a, b]$ has exact PR area between*

$$L^{\mathrm{PR}}_{a,b} = \int_0^1 \min_{\ell(z) \leq j \leq u(z)} \mathrm{Prec}_{a,j}\, dz, \tag{10}$$

$$U^{\mathrm{PR}}_{a,b} = \int_0^1 \max_{\ell(z) \leq j \leq u(z)} \mathrm{Prec}_{b,j}\, dz. \tag{11}$$

Both bounds can be computed in $O(G)$ time and space.

Proof. The quantile representation of $\sum_j q_j \mathrm{Prec}_j$ integrates the precision at the quantile index over z. That index lies in the displayed window, and its precision is between the corresponding endpoint precisions. Taking the window minimum and maximum and integrating yields the bounds. Breakpoints are the union of the two endpoint recall sequences. Merge them in linear time. Both ends of the index window move only right; standard monotone deques maintain the precision minimum and maximum in amortized constant work per inserted or removed index. Each index is inserted and removed at most once. □

These are extrema over a relaxed p-box model, not necessarily values attained by an intermediate VUS width. This relaxation permits a valid bound despite oscillating precision. We claim the VUS-specific coordinate order, merge partition, and adaptive use of these bounds, not the general p-box expectation construction.

3.3 Adaptive evaluation and decisions

Evaluate the endpoints of each merge-free interval. If an active interval $[a, b]$ has $r = b - a - 1$ unevaluated interior widths, add $rL_{a,b}$ and $rU_{a,b}$ to the lower and upper sums, respectively. Add evaluated scalar values exactly once, then divide by $L + 1$. This encloses the target discrete average; it does not replace it with a continuous-width integral.

Refine the interval with the largest contribution to uncertainty by evaluating its midpoint and replacing it by two subintervals. Stop once both metric intervals have half-width at most ε, or when every width is evaluated. The returned estimate is the midpoint. One model outranks another whenever its lower bound exceeds the other's upper bound; overlapping intervals produce no ordering assertion.

With m distinct evaluated widths, preparation, endpoint calculation, linear p-box bounds and heap refinement take $O(n \log(Q+1) + mQ + m \log m)$ time. A fixed-size endpoint-curve cache retains at most eight curves; evicted endpoints can be recomputed a constant number of times per refinement. Space is $O(n + Q + m)$, with no width-by-support matrix. The artifact records both distinct widths and actual curve evaluations. The bound is output-sensitive to the requested precision and geometry; neither $m \ll L$ nor a wall-clock improvement is guaranteed.

The theorems are real-arithmetic statements. The implementation uses float64, clips ulp-scale recall drift, and adds an empirical 10^{-12} guard to adaptive endpoints. It is not a directed-rounding interval library. Experiments independently audit enclosure within 10^{-10} and use a $2 \cdot 10^{-10}$ separation margin for numerical ordering checks. These allowances are reported separately from the requested approximation half-width. Setting $\varepsilon = 0$ evaluates every width.

4 INCREMENTAL PR MODEL SELECTION

Fixed-tolerance evaluation asks each model to meet an absolute error target before using the resulting intervals to select a winner. This can over-refine already separable models and repeat preparation during exact fallback. We therefore also study an incremental PR race. It was designed after inspecting the standalone fixed-tolerance width-refinement results; its measurements are a separate post-hoc follow-up on the same 126 series/width selection tasks.

For each of M models, prepare its sparse evaluator once and initialize the endpoints of every merge-free width block. Retain all observed PR scalars, an endpoint-curve cache, and a heap of unevaluated interiors. An interior (a, b) has priority $(b - a - 1)(U^{\mathrm{PR}}_{a,b} - L^{\mathrm{PR}}_{a,b})$. Its midpoint refinement reuses preparation and updates the model's global PR enclosure $[\ell_k, u_k]$. The reused endpoint oracle still computes ROC, but only PR ordering controls selection; there is no predetermined scalar-error target.

After initialization and every refinement, rescreen *all* models. Return a singleton k only if $\ell_k > u_j + \eta$ for every $j \neq k$, where $\eta = 3 \cdot 10^{-10}$. Otherwise the current candidate set is

$$\mathcal{C} = \{k : u_k \geq \max_j \ell_j - \eta\}.$$

This is repeated screening, not permanent deletion; a model outside $\mathcal{C}$ retains its state and is checked again. Among nonexact candidates, refine the model with the widest global PR interval, breaking ties by model name, and bisect its highest-priority width interior. If all candidates are exact, return every candidate whose exact PR is within $\tau = 10^{-10}$ of their maximum. Exact means use already observed widths; no second preparation or separate exact evaluator is called. The no-pruning ablation keeps all models in $\mathcal{C}$ whenever the singleton condition fails, retaining the same stopping and refinement rules.

Proposition 1 (Finite decision resolution). *With valid real-arithmetic width enclosures and $\eta > \tau$, the race returns the exact τ-tie set and evaluates at most $M(L + 1)$ distinct model/width pairs.*

Proof. A separated winner exceeds every other true PR by more than η. A currently excluded model has true PR more than

η below some model's lower bound, hence cannot belong to the global τ-tie set. At least one candidate exists. Each refinement evaluates one unseen width in a nonempty interior; active interiors are disjoint. Therefore at most $M(L+1)$ distinct pairs can be observed. If separation does not occur, all remaining candidates eventually become exact and the tie rule resolves them. □

Let T be the number of refinements and m_k the number of distinct widths queried for model k. In addition to one preparation per model, endpoint work is $O(Q\sum_k m_k)$, heap work is $O(T\log(L+1))$, and rescreening is $O(MT)$. With the default cache, actual curve calls are at most the initialization count plus $3T$; both actual calls and distinct widths are recorded. Close scores can require full evaluation, and decision bookkeeping can outweigh avoided curves. The float64 implementation retains the empirical 10^{-12} guard; the ordering margin is a numerical allowance, not a directed-rounding proof.

5 PRIOR ART AND IMPLEMENTATION COMPARISONS

Boniol et al. introduce optimized VUS implementations with static/dynamic label splitting, maximum-buffer support restriction, and memory tradeoffs [1]; these are substantive prior optimizations. Klein's Apache-2.0 fast-VUS implementation, first committed on 12 October 2026, computes support-score cumulative sums, globally sorts all scores, and evaluates the selected thresholds at every width [2]. At audited revision `1bae859f`, passing `thre=n` gives full-threshold PR with the target's dynamic mass, sum-and-clip ramps, expanded-interval existence reward, and right-endpoint integration. Its API returns PR only; our measured API computes ROC and PR together. Default 250-threshold and full-threshold calls are distinguished in every comparison. We do not claim support restriction, sorting support scores, or cumulative sums as new. The remaining exact reduction removes inactive threshold events, including their ROC contribution, and obtains prediction counts without globally sorting all scores. Width enclosures are a separate contribution.

Sylligardos reports preliminary VUS speedups up to two orders on CPU and three on GPU over approximately 1,900 TSB-UAD series [8]. That workshop paper does not disclose the detailed algorithm, and a corresponding released implementation was not located in our 20 October 2026 source audit. Its mechanism could overlap with ours. We therefore do not claim the first fast VUS implementation, the first 100-fold acceleration, or demonstrated superiority over that method.

The released Go `adscore` package implements an incremental dense rank sweep [7]. Native execution of its unmodified v0.0.23 package confirms different semantics: static positive mass, original-interval existence, maximum-combined overlapping ramps, and trapezoidal PR. On $y=(0,1,0,0,1,0)$, $s=(1,2,3,4,5,6)$, and $L=4$, its ROC/PR values are $(0.635984, 0.638587)$, versus $(0.689265, 0.683409)$ for the pinned target. It is prior art for rank sweeps, but not an interchangeable exact-target runtime baseline.

Current dtaianomaly commit `557b4357` provides a compiled, parallel support-scan implementation with the target semantics [5]. We execute its original method bodies and compiled kernel unchanged through a dependency-isolating adapter, obtain both integrals from one curve calculation, and use one thread. Its support restriction differs from eliminating entire negative-only threshold runs and aggregating all support events in one pass. Both default 250-threshold and full-threshold variants are evaluated.

P-box lower and upper expectations for nonmonotone functions are established by Utkin and Destercke [10]; Troffaes and Destercke develop p-box inference on totally preordered spaces [9]. Our application uses the ordered support-event index as the state space and endpoint recall as CDF bounds. The VUS-specific work is the merge partition, coordinate-order proof, and adaptive integration over integer widths. The source audit located no disclosed VUS algorithm with these width enclosures, but this bounded search does not establish priority.

6 EXPERIMENTAL PROTOCOL

The protocol, exact source hashes, score archive selection, and a pre-outcome addendum are included with the artifact. The dtaianomaly comparator was selected before execution; a separate, retrospectively written addendum records that decision and its timing. One thousand independently seeded random cases vary length, label density, ties, constant scores, and buffer width; six structured cases cover edges and overlapping/adjacent expansions. Separately authored validation scripts add 1,960 exhaustive tiny cases and 120 new random cases, including extreme finite scores and arbitrary boundary labels.

We use three released score archives from Zhu et al. [11]: Chronos-Bolt base, Time-MoE 200M, and TimesFM 1.0 200M, each at context length 64. Models were chosen by a frozen family/context/lexicographic rule. All common valid series from the 350-series TSB-AD-U evaluation split [4] are included; no model is trained or queried. Scores are read unchanged. Their source README documents leading zero padding for Chronos-Bolt and Time-MoE; we do not reproduce that paper's corrected model leaderboard. Original dataset licenses vary, and raw data and scores are not redistributed.

The main experiment evaluates maximum widths 10 and 30. Each timing includes validation, preprocessing, sorting or histogramming, and integration. Three repetitions use a seeded randomized method order; JIT compilation/first-call time and downloads are excluded and recorded separately. The machine is an Apple M1 Pro, Python 3.12.4, with all compute libraries restricted to one thread. We compare NumPy sparse, Numba sparse, NumPy dense, Numba dense with precomputed nearest distances, and a separately implemented Numba dense method constructing every original sum-and-clip ramp directly. The principal speed ratio uses the fastest median among all three dense implementations for each input, favoring the dense baselines.

The final timed output for each full trace is cross-checked between all five implementations. Official full-threshold and sampled-threshold equivalence are checked on fixed 256-position crops around each first anomaly, when both classes occur. Published external runtimes are measured on the prespecified first 12 common series (all NAB, lengths 1,000–10,319), with all three score vectors and both widths. Official 250-threshold timing has one repetition; dtaianomaly has three. The five in-repository

methods use randomized order. External methods use the fixed order official-250, dtaianomaly-250, then dtaianomaly-exact. This external timing subset is not representative of all dataset families.

Two supplementary protocols were fixed after review of the main setup and before their respective executions. A mechanism ablation replaces only the score-count preprocessing with a global full-score sort, then calls the identical compiled sparse integration kernel. It crosses two lengths ($10^4, 10^5$), four positive densities ($0.001, 0.01, 0.1, 0.5$), four interval counts ($1, 10, 100, 1000$), and four maximum widths ($0, 10, 30, 100$). Sixteen geometrically infeasible cells are excluded before execution, leaving 112 cases. Five implementations run in randomized order for three repetitions. The second supplement tests the actual published implementation across source families; its selection and censoring rules are reported with the results.

For sampling effects, all three unordered model pairs per series/width are examined. A reversal requires opposite signs and absolute metric differences above 10^{-10} in both calculations; ties are reported separately. Series, not score vectors or width/model repetitions, are the units for uncertainty summaries. These are changes in ordering induced by metric computation, not claims that one model is superior.

The padding sensitivity protocol was fixed after observing the original results and padding inventory, before evaluating its transformations. For Chronos-Bolt and Time-MoE, *first-score fill* replaces the first 64 scores with the score at index 64, retaining full labels and all TimesFM scores. *Common suffix* evaluates positions 64 onward for every model. Both interventions use the same 350 series, model pairs, and widths. The suffix changes prevalence, threshold ranks, and boundary geometry; 21 series contain positive labels in the removed prefix. Neither intervention reconstructs every preprocessing choice in the source model paper.

7 RESULTS

7.1 Exact evaluation against released fast-VUS on two machines

A new complete run compares the unchanged fast-VUS module at full thresholds against Sparse VUS and the three dense implementations. Each machine evaluates all 2,100 original model/width cases on 350 series from 23 families. Each call includes validation, preprocessing, and integration; three repetitions use randomized method order, one compute thread, and excluded JIT warmup. The M1 Pro run uses Python 3.12.4, NumPy 1.26.4, and Numba 0.59.1. The Linux hosted runner identifies an AMD EPYC 9V74 and uses Python 3.12.15, NumPy 2.5.4, and Numba 0.68.0. Ordinary OS and hosted-hardware background load is uncontrolled. The Linux run is a separate execution, not an external laboratory reproduction.

Table 1 reports within-machine ratios. Fast-VUS returns only PR; Sparse VUS computes both ROC and PR. The per-case fastest-tested comparator is the minimum median runtime among fast-VUS and all three dense evaluators, excluding Sparse VUS itself. This is an oracle comparison, not a deployable selector. The maximum PR difference from fast-VUS is $2.78 \cdot 10^{-15}$; the largest discrepancy across all exact methods is $3.70 \cdot 10^{-13}$.

Fast-VUS ratios range from 1.75 to 46.45 on M1 Pro and 1.52 to 17.97 on Linux. Against the fastest-tested comparator, ranges are 0.40–7.75 and 0.36–8.40. Whole-family bootstrap intervals for the latter geometric means are [1.53,2.42] and [1.11,1.68], respectively. The corpus and machines support a bounded measured advantage, with substantial dense-baseline slowdown cases; they do not settle the unavailable Sylligardos comparison. The sums of isolated per-case medians are 14.78 versus 104.52 seconds for Sparse VUS and fast-VUS on M1 Pro, and 14.39 versus 55.14 seconds on Linux. These sums are not measured batch or detector-pipeline wall times.

7.2 Error bounds and the cost of width refinement

The width-refinement experiment retains the metadata-selected 42 series covering all 23 families and all three released score vectors. Maximum widths 30, 100, and 300 give 378 model/width cases; the width-300 setting is an explicit stress condition. Both $\varepsilon = 0.01$ and 0.001 bound the half-widths of both returned metric intervals. The M1 Pro timing protocol again includes whole-call preprocessing and uses three randomized-order repetitions. Table 2 compares against full-width Sparse VUS on these same inputs.

All 378 cases at both tolerances contain the full-width ROC and PR values within the separately declared 10^{-10} numerical audit allowance. Maximum midpoint errors are 0.003260 and 0.000713 for tolerances 0.01 and 0.001. However, fewer width evaluations do not generally imply faster execution. Across all three maximum widths, median runtime ratios are 0.966 and 0.467, with adaptive evaluation faster in 173 and 40 of 378 cases, respectively. At tolerance 0.01 and $L = 300$, the geometric mean ratio is 1.370 (series-bootstrap interval [1.207,1.557]); at the tighter tolerance, all three width groups have geometric mean ratios below one. The present implementation therefore provides useful enclosures, but the standalone fixed-tolerance strategy is often slower than exact evaluation.

As a numerical decision replay, nonoverlapping PR intervals identify unique winners in 21, 17, and 15 of the 42 series at tolerance 0.01 for widths 30, 100, and 300; the corresponding counts at tolerance 0.001 are 33, 32, and 34. All asserted winners agree with exact PR using the declared $2 \cdot 10^{-10}$ separation margin. These replay counts do not include additional decision-workload timing or establish a model-selection speedup.

Separately authored validation directly constructs full-threshold curves for every both-class binary label vector of lengths two through seven, using four seeded score patterns per vector. It checks 13,776 widths and width enclosures and 960 discrete averages. The largest curve discrepancy is $6.67 \cdot 10^{-16}$ and the largest enclosure violation is $3.34 \cdot 10^{-16}$, within the declared $2 \cdot 10^{-12}$ allowance. A further 760 linear-programming comparisons of p-box extrema have maximum discrepancy $3.34 \cdot 10^{-16}$. Exhaustiveness concerns labels only, not all rankings or tie patterns; these checks audit the implementation and complement the real-arithmetic proofs.

7.3 Complete model-selection workloads

We select the maximum exact PR among all three released score vectors, retaining every model within 10^{-10} of the maximum. Both experiments use the same fixed 42-series metadata selection and $L \in \{30, 100, 300\}$ (126 tasks), with no case exclusions. Every

Machine	Comparator	Median	Geometric mean (95% interval)	Faster	Slower or equal
M1 Pro	fast-VUS	8.38	8.69 [8.32, 9.05]	2100	0
M1 Pro	Fastest tested	2.04	1.98 [1.89, 2.07]	1891	209
EPYC 9V74	fast-VUS	4.55	4.50 [4.34, 4.68]	2100	0
EPYC 9V74	Fastest tested	1.35	1.40 [1.34, 1.46]	1526	574

Table 1: Comparator runtime divided by Sparse VUS runtime; values above one favor Sparse VUS. Every row covers 2,100 cases. Intervals resample whole series and describe geometric means on the fixed corpus. Both methods use all thresholds; fast-VUS computes PR only.

ε	L	Median ratio	Faster / 126	Distinct widths (%)	Curve calls (%)
0.01	30	0.835	27	16.1	16.1
0.01	100	0.905	46	7.9	7.9
0.01	300	1.313	100	4.7	7.3
0.001	30	0.615	4	54.8	80.6
0.001	100	0.426	11	39.6	74.3
0.001	300	0.391	25	25.4	53.8

Table 2: Full-width exact runtime divided by adaptive runtime. Width and curve-call columns are medians relative to $L + 1$; cache eviction can make curve calls exceed the number of distinct widths. Values below one indicate that adaptive evaluation is slower.

timed call includes all model preparation and all subsequent refinement or fallback; score loading and warmed JIT compilation are excluded. Model training and inference are not measured. Width 300 remains an explicit stress regime.

Fixed tolerance with exact fallback. At initial half-width 0.01, the selector's median speed ratios against sparse exact selection are 0.579×, 0.569×, and 0.738×. All 126 decisions agree, but the complete selector is slower at every width in the median. Pruning reduces exact fallback model evaluations from 63/75/81 to 53/63/69, without recovering this overhead. The fixed-250 comparator changes the selected set in 4/5/5 tasks. Of these, 3/4/4 have a unique exact winner; the remaining case at each width breaks an exact three-way tie.

Incremental PR racing. The follow-up removes the fixed accuracy target and reuses each model's prepared state (Section 4). This design was chosen after observing the fixed-tolerance overhead; it is not a preregistered independent generalization study. Fresh timings on both architectures retain the same inputs and separate each executable exact comparator. Every racing decision agrees with exact selection on both platforms.

Median gain versus total cost. At width 300 on macOS, summing the 42 case medians gives 10.615 s for racing versus 6.249 s for sparse exact evaluation (more total time). At width 300 on Linux, summing the 42 case medians gives 10.522 s for racing versus 7.051 s for sparse exact evaluation (more total time). Thus the distribution median and the total workload cost must be read separately; a positive median speed ratio does not establish net time savings. The no-pruning ablation retains the separated-winner rule and refines all models when no winner separates. Its median time divided by the pruning variant's time, at the three increasing widths, is macOS: 1.022, 1.032, 1.094; Linux: 1.018, 1.056, 1.088; these values isolate the incremental effect of pruning in this workload. All bound-separation decisions are numerical audits under the stated 10^{-10} enclosure allowance, not formal Float64 certificates. The artifact includes absolute totals, series-clustered intervals, full-evaluation counts, and every slower task.

7.4 Corpus-wide sampling effects and supporting measurements

The original fixed-corpus run provides the following threshold-sampling analysis and complementary dense-baseline measurements. These timings are retained as a separate run; they are not pooled with the current published-code comparisons. The completed experiment contains 350 series from 23 source families, 1050 released score vectors, and 2100 score/width evaluations. Sequence lengths range from 1,000 to 900,000. The independent final artifact gate passes all coverage, finite-value, source-hash, internal, crop, and external-comparator checks. The maximum absolute numerical discrepancy is 3.69×10^{-13}.

The median compiled-sparse speedup over the fastest dense baseline is 2.06×, with geometric mean 2.02×. The 95% series-cluster bootstrap interval for the geometric mean is [1.93, 2.12]; resampling whole source families gives [1.56, 2.48]. These are descriptive uncertainty summaries for the fixed corpus. Sparse evaluation is faster in 1916 of 2100 cases, while 184 cases are slower or equal. The full speedup range is 0.27–7.50×. Figure 2 retains these slowdowns rather than filtering dense support.

For VUS-ROC, strict reversals occur on 2 unique series across 2 families, and the unique top model changes in 2 of 700 series/width

Platform	L	vs. sparse	vs. oracle	wins/42	exact ms	race ms	widths
macOS	30	0.829	0.829	8	9.08	12.96	18.3%
macOS	100	0.715	0.678	11	16.57	23.72	16.3%
macOS	300	1.177	1.063	22	38.34	41.37	10.5%
Linux	30	0.817	0.796	8	9.11	13.93	18.3%
Linux	100	0.757	0.704	10	18.90	27.04	16.3%
Linux	300	1.376	1.216	26	49.06	45.16	10.5%

Table 3: Incremental PR selection, 42 tasks per row. Ratios above one favor racing. The oracle chooses the fastest of three separately measured compiled exact methods per task; it is not deployable. Wins compare racing to that oracle; latency columns are medians in milliseconds for sparse exact and racing. Widths is the median fraction of the $3(L + 1)$ distinct model-width queries evaluated. Recomputed endpoint curves are separately counted in the released records.

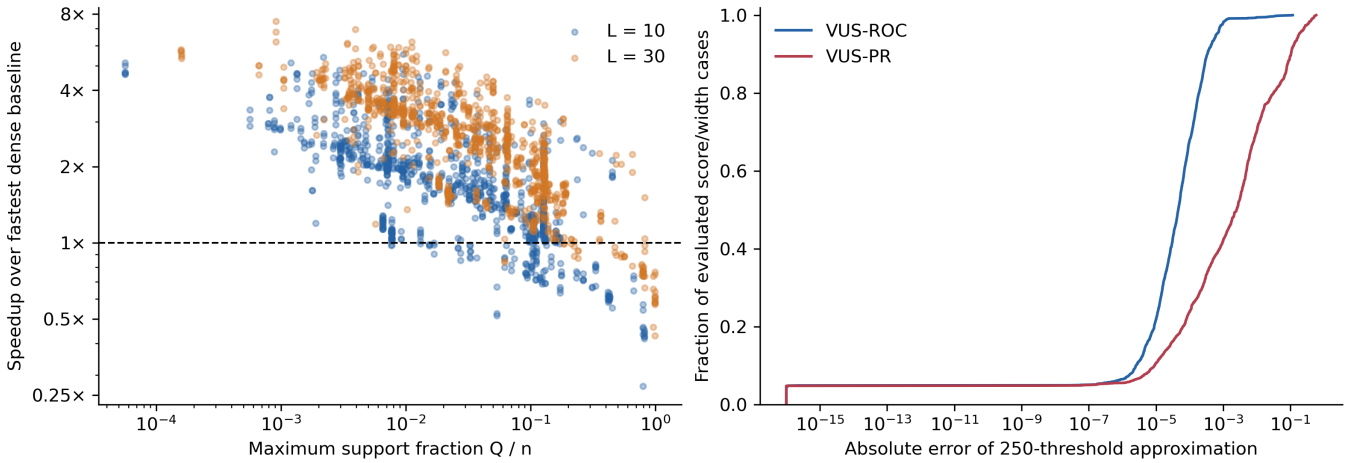


Figure 2: Left: all 2100 paired runtime comparisons, using the fastest of three dense implementations per input. Right: threshold-sampling error on the same fixed score vectors. Zero errors are displayed at 10^{-16} for the logarithmic axis.

Metric	Median error	95th percentile	Maximum error	Reversals / pairs
VUS-ROC	4.59×10^{-5}	6.21×10^{-4}	0.117	5 / 2100
VUS-PR	2.50×10^{-3}	0.178	0.583	84 / 2100

Table 4: Errors are absolute differences between the same metric at 250 sampled thresholds and all thresholds. Reversals require strict opposite order in both calculations.

cases. Exact and sampled comparisons contain 96 and 95 ties, respectively.

For VUS-PR, strict reversals occur on 55 unique series across 13 families, and the unique top model changes in 50 of 700 series/width cases. Exact and sampled comparisons contain 96 and 95 ties, respectively.

The largest observed PR discrepancy occurs on the 195,000-point UCR HumanActivity series with Time-MoE scores and $L = 30$: exact VUS-PR is 0.666890, versus 0.084159 at 250 thresholds. The scores are unchanged between calculations.

On the bounded 12-series NAB external comparison, the median speed ratios relative to compiled sparse evaluation are 53.86× for full-threshold dtaianomaly, 8.92× for 250-threshold dtaianomaly, and 149.29× for official 250-threshold evaluation. The latter two compare exact computation with a coarser approximation. These results do not establish the same advantage on other families or against the unavailable accelerated VUS implementation.

In the frozen fixed-support scaling experiment ($Q = 102$, $L = 30$), the one-million-point case takes 0.0626 seconds for compiled sparse evaluation, an 8.92× speedup over the fastest dense implementation, with five randomized-order repetitions.

7.5 Mechanism and the density/width tradeoff

On the 112-case synthetic grid, replacing full-score sorting with support-score counting gives a geometric-mean end-to-end speedup of 1.11× with the same integration kernel, range 0.94–1.80×. The adapter deliberately holds support grouping and all

other preparation fixed, isolating the score-counting choice. Relative to the fastest dense implementation, the complete sparse method has geometric-mean speedup 1.37×, range 0.26–7.53×. There are 69 cases with at least a 1.1× speed ratio and 34 with ratio at most $1/1.1$. The remaining cases lie between those thresholds. Thus the count-preprocessing improvement is modest on this grid; sparse event integration supplies the larger opportunity as support shrinks and the width range grows.

7.6 Published implementation across source families

The post-review supplement selects the shortest and median-length series within each family, deduplicating coincident choices. It contains 42 full-length series from all 23 families, uses Chronos-Bolt base scores and $L = 30$, and preserves all vectors. Methods run in a seeded random order, with three repetitions per method. Each call has a ten-second post-warmup time limit; late or unfinished calls are recorded as right-censored. Compilation and subprocess startup are excluded. Table 5 reports speed ratios only for completed pairs, whose subset differs across threshold modes. No numerical speed ratio is imputed for a censored result.

The maximum absolute discrepancy for returned outputs, compared with the corresponding exact or 250-threshold main result, is 9.55×10^{-15}. All selection, timeout, environment, and per-call receipts are retained.

7.7 Padding sensitivity and effect sizes

Table 6 reports all pairs, including small effects. All 350 series and 23 source families remain eligible in each variant. The largest newly computed exact-versus-dense discrepancy is $2.37 \cdot 10^{-13}$; 144 transformed small-case upstream checks also pass. Of 84 original PR reversals, 75 persist after filling and 74 on the suffix, all in the same direction. Filling removes nine and introduces 18; cropping removes ten and introduces 21. Thus documented zero padding alone does not explain the effect, while individual changed orders depend on preprocessing.

Most reversals have small margins. Requiring both exact and sampled PR gaps to be at least 10^{-3} leaves 38, 41, and 44 reversals in the three variants; at 10^{-2} each has five. On the common suffix the median, 95th-percentile, and maximum absolute PR errors are 0.002489, 0.177261, and 0.582639. Giving source families equal weight yields reversal prevalences of 2.44%, 5.45%, and 4.31%, with descriptive family-bootstrap intervals [1.12,4.06]%, [2.12,10.32]%, and [1.84,7.59]%, respectively. These differ from the pooled pair fractions 4.00%, 4.43%, and 4.52%; singleton families cannot support precise population estimates. The result concerns approximation-induced metric order, not a change in detector quality or a claim that the exact metric is the best model selector.

8 LIMITATIONS AND REPRODUCIBILITY

The target is a specific, pinned scalar implementation. Exactness does not resolve scientific criticisms of the underlying metric, extend it to all historical variants, or produce the full surface. The acceleration depends on sparse expanded support and can disappear for dense support. Adaptive bounds can require all widths and add overhead; tolerance is not a guaranteed speedup. Real-arithmetic enclosures and float64 checks are distinct: no machine-level interval certificate or unconditional numerical ordering guarantee is claimed. Floating-point equivalence is numerical, not bitwise. The NumPy/Numba dense implementations are independently written algorithmic baselines, not published competing systems. Published-code comparisons retain API and threshold-resolution differences. The actual released dtaianomaly comparison is bounded, and the Sylligardos fast implementation remains unbenchmarked.

The repository includes source, tests, a rigorous derivation, frozen protocols, download and source-checksum receipts, per-input timing and metric records, negative results, and scripts for regenerating figures and this manuscript. All experiments reuse released score vectors on CPUs; no models are retrained. Experiments were executed on 19–20 October 2026 UTC. Development and in-session code review used AI assistance; the separately authored checks are distinct test implementations within this session, not an external laboratory reproduction.

9 CONCLUSION

Inactive threshold events can be integrated analytically, giving exact scalar VUS with work governed by expanded anomaly support. On the tested corpus this evaluator is faster than released fast-VUS on both CPU systems, while strong dense methods retain favorable cases. Merge-aware coordinate order and p-box bounds additionally support width enclosures and incremental PR decisions. These extensions reduce distinct width queries but introduce substantial overhead: even where racing improves the median speed ratio, the sum of per-task medians increases. They remain optional APIs; full-width exact evaluation is the default.

A A FIXED UNIFORM-RANK GRID CAN LOSE ALMOST ALL PR AREA

The following proposition concerns the particular uniform-rank sampling rule used by the pinned reference, with all input scores distinct. It does not apply to every adaptive threshold-selection scheme.

Proposition 2. *Fix an integer threshold budget $T \geq 2$, a maximum width $L \geq 0$, and $\varepsilon > 0$. There are finite-score binary-label sequences for which the absolute difference between exact VUS-PR and the reference's T-threshold VUS-PR is greater than $1 - \varepsilon$.*

Proof. Take one anomaly interval of length m, with at least $\lfloor L/2 \rfloor$ positions of padding on both sides. Give its original positives score ranks $2, \dots, m+1$, put the highest-scored point outside the maximum-width support, and give all buffer positions the next consecutive ranks. Give all other positions distinct lower scores. Append further lower-scored negatives at the right, preserving the existing geometry and scores.

For every $w \leq L$, the jth original positive is admitted before any buffer point. Its recall is j/m, precision is $j/(j+1)$, and existence reward is one. Recall reaches one after the last original positive. All subsequent PR increments are zero, so exact VUS-PR is

$$E_m = \frac{1}{m} \sum_{j=1}^{m} \frac{j}{j+1} = 1 - \frac{H_{m+1} - 1}{m}, \tag{12}$$

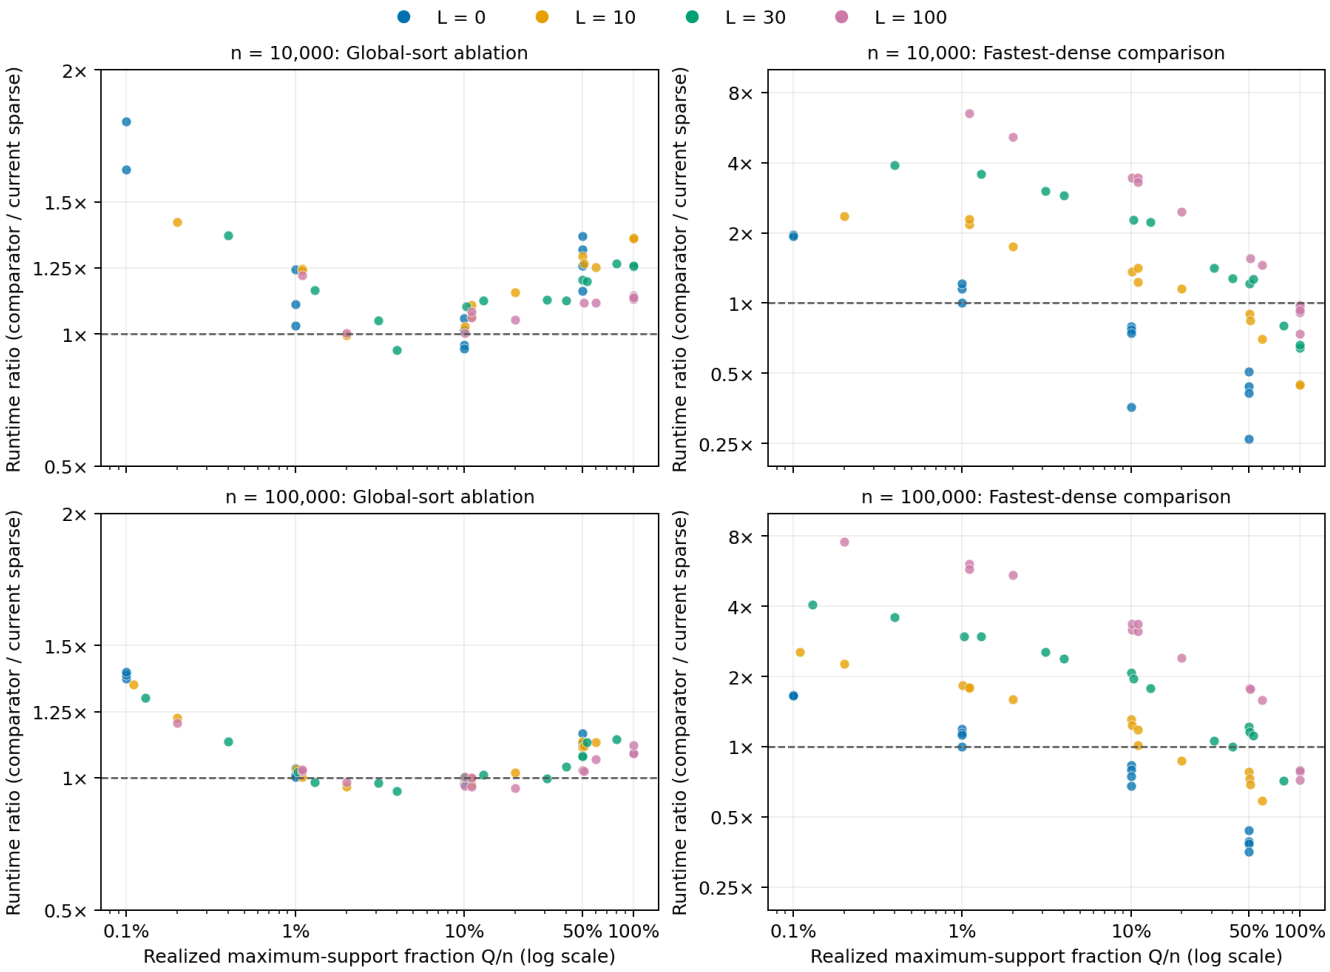


Figure 3: Mechanism ablation: each point is one predeclared synthetic cell, using medians of three randomized repetitions. Support fraction Q/n is measured after expansion. The dense comparator is the fastest of three dense implementations. The full-sort variant keeps support grouping and the compiled sparse integration kernel fixed.

dtaianomaly mode	Completed	Censored	Families represented	Median ratio
250 thresholds	42	0	23	8.17×
All thresholds	35	7	19	143.85×

Table 5: Actual pinned dtaianomaly code on 42 source-stratified full-length inputs. Ratios divide the published baseline's median time by compiled sparse exact evaluation. The 250-threshold row compares an approximation with exact computation. Completion-conditioned medians are not full-corpus runtime estimates. Family counts require at least one completed pair within the family.

where H_k is the kth harmonic number.

The sampled grid has one-based ranks $1 + \lfloor j(n-1)/(T-1) \rfloor$, for $j = 0, \ldots, T-1$. For sufficiently large n, its second rank $k_n = 1 + \lfloor (n-1)/(T-1) \rfloor$ contains all maximum-width support. The first sampled recall is zero. Let $A_w = \sum_i a_i^{(w)}$ denote the full label mass. At the second threshold, $\mathrm{TP} = A_w$, $B = A_w - m$, and $P = (m + A_w)/2$. Thus clipped recall and existence reward both equal one, and precision is A_w/k_n. No later sampled threshold adds PR area. The sampled VUS-PR therefore equals $\overline{A}/k_n$, where $\overline{A} = (L+1)^{-1} \sum_{w=0}^{L} A_w \le m + L$.

First choose m so that $(H_{m+1} - 1)/m < \varepsilon/2$; then append enough negatives that k_n exceeds all support ranks and $2(m+L)/\varepsilon$. The exact area exceeds $1 - \varepsilon/2$ and the sampled area is less than $\varepsilon/2$, proving the claim. □

Variant	PR reversals	Series	Top changes	ROC reversals
Released	84	55	50	5
First-score fill	93	59	52	12
Common suffix	95	61	48	8

Table 6: Each row has 2,100 unordered model-pair comparisons and 700 series/width decisions. A reversal has opposite signs with both gaps above 10^{-10}. Top changes require a unique best model in both calculations.

The artifact includes deterministic checks at $T = 250$, $m \in \{1, 10\}$, $L \in \{0, 6\}$, and $n \in \{10^4, 10^5\}$ against the two closed-form values.

REFERENCES

[1] Paul Boniol, Ashwin K. Krishna, Marine Bruel, Qinghua Liu, Mingyi Huang, Themis Palpanas, Ruey S. Tsay, Aaron Elmore, Michael J. Franklin, and John Paparrizos. 2025. VUS: Effective and Efficient Accuracy Measures for Time-Series Anomaly Detection. *The VLDB Journal* 34, 3 (2025), 1–32. https://arxiv.org/abs/2502.13318

[2] Tom Louis Klein. 2026. fast-VUS. https://github.com/tomlouisklein/fast-VUS Software, first commit 12 October 2026; audited revision 1bae859f, 16 October 2026.

[3] Krzysztof Kotowski, Christoph Haskamp, Jacek Andrzejewski, Bogdan Ruszczak, Jakub Nalepa, Daniel Lakey, Peter Collins, Aybike Kolmas, Mauro Bartesaghi, Jose Martinez-Heras, and Gabriele De Canio. 2024. European Space Agency Benchmark for Anomaly Detection in Satellite Telemetry. https://arxiv.org/abs/2406.17826 arXiv:2406.17826, revised 17 August 2025.

[4] Qinghua Liu and John Paparrizos. 2024. The Elephant in the Room: Towards a Reliable Time-Series Anomaly Detection Benchmark. In *Advances in Neural Information Processing Systems*, Vol. 37. 108231–108261. https://github.com/TheDatumOrg/TSB-AD

[5] ML-KULeuven. 2026. dtaianomaly. https://github.com/ML-KULeuven/dtaianomaly Software, audited revision 557b4357, 13 October 2026.

[6] John Paparrizos, Paul Boniol, Themis Palpanas, Ruey S. Tsay, Aaron Elmore, and Michael J. Franklin. 2022. Volume Under the Surface: A New Accuracy Evaluation Measure for Time-Series Anomaly Detection. *Proceedings of the VLDB Endowment* 15, 11 (2022), 2774–2787. https://doi.org/10.14778/3551793.3551830

[7] P. Stergiotis. 2026. boxer/adscore. https://pkg.go.dev/github.com/stergiotis/boxer/public/analytics/timeseries/adscore Software version 0.0.23.

[8] Emmanouil Sylligardos. 2025. Fast or Accurate? Rethinking Time Series Anomaly Detection. In *VLDB Ph.D. Workshop*. https://www.vldb.org/2025/Workshops/VLDB-Workshops-2025/PhD/PhD25_17.pdf Section 4.

[9] Matthias C. M. Troffaes and Sebastien Destercke. 2011. Probability Boxes on Totally Preordered Spaces for Multivariate Modelling. *International Journal of Approximate Reasoning* 52 (2011), 767–791. https://doi.org/10.1016/j.ijar.2011.02.001

[10] L. Utkin and S. Destercke. 2009. Computing Expectations with Continuous P-Boxes: Univariate Case. *International Journal of Approximate Reasoning* 50, 5 (2009), 778–798. https://doi.org/10.1016/j.ijar.2009.02.004

[11] Xiaokun Zhu, Louis Carpentier, and Mathias Verbeke. 2026. When Foundation Models are One-Liners: Limitations and Future Directions for Time Series Anomaly Detection. In *International Conference on Learning Representations*. https://openreview.net/forum?id=H27kvyG4qf Released scores: KU Leuven RDR, V1, doi:10.48804/MCVI4N.